
Covariance-Driven Graph Embedding for Real-Time Traffic State Prediction

Sai Vamsi Aliseti
Department of Computer Science
University of California, Santa Barbara
saivamsi@ucsb.edu

Vikas Kalagi
Department of Computer Science
University of California, Santa Barbara
vikaskalagi@ucsb.edu

Abstract

1 Accurate traffic forecasting is essential for enhancing urban mobility and enabling
2 intelligent transportation systems. This paper introduces a novel approach to traffic
3 prediction by leveraging temporal graph structures to capture spatial and temporal
4 dependencies within traffic networks. Each graph, defined by its nodes, edges,
5 and adjacency matrix, represents the traffic conditions at a specific time step. Our
6 method employs a feature-level analysis by computing pairwise covariances be-
7 tween nodes in the current graph and those in preceding time steps. To extract
8 meaningful patterns, we perform eigenvalue decomposition to determine the prin-
9 cipal eigenvector. Node features are then projected onto this eigenvector, and
10 cosine similarity is utilized to identify the most relevant nodes for each feature.
11 This process ensures a robust and adaptive feature selection mechanism, enabling
12 our approach to effectively capture complex traffic dynamics for more accurate
13 predictions. Please find the code here - GitHub.

14 1 INTRODUCTION

15 In the modern era of interconnected systems, traffic prediction plays a pivotal role in optimizing trans-
16 portation networks, reducing congestion, and enabling smart city applications. Accurate modeling of
17 traffic dynamics is inherently challenging due to the spatial and temporal dependencies among traffic
18 nodes, as well as the influence of external factors. Graph-based models provide a powerful framework
19 to capture these dependencies, where nodes represent traffic points, edges denote connections, and
20 node features encapsulate traffic-related attributes such as speed, flow, and density. In this study, we
21 propose a novel approach to model traffic dynamics by leveraging graph structures with adjacency
22 matrices, where each graph represents a snapshot of the traffic network at a given time instance.

23 Our methodology focuses on incorporating temporal graphs, $\mathcal{G}^{(t)} = (V, E, \mathbf{A})$, where each node
24 is associated with F features and T time instances. To enhance predictive accuracy, we compute
25 pairwise covariance matrices between nodes in the current graph $\mathcal{G}^{(t+1)}$ and preceding graphs $\mathcal{G}^{(t)}$,
26 perform eigenvalue decomposition, and utilize the principal eigenvectors to select the most relevant
27 nodes based on cosine similarity. This selection process identifies key spatial relationships while
28 dynamically adapting to temporal patterns. Additionally, a learnable attention mechanism refines
29 the spatial weights by accounting for higher-order neighborhood effects, ensuring that our model
30 captures both local and global correlations. The aggregated node features across S time steps are
31 further processed through a multi-layer framework to compute the final embeddings.

32 By combining graph theory, statistical methods, and attention-based learning, our framework aims to
33 provide an interpretable and robust mechanism for traffic forecasting. The results demonstrate its
34 potential to enhance decision-making for traffic management and urban planning.

2 LITERATURE REVIEW

GraphSAGE-based Dynamic Spatial–Temporal Graph Convolutional Networks (DST-GCN) [4] have recently gained significant attention in the domain of traffic prediction, primarily due to their ability to capture both spatial dependencies among traffic nodes (e.g., road intersections or sensors) and temporal patterns in traffic flow data. The GraphSAGE [2] framework, which utilizes inductive graph neural networks, aggregates information from neighboring nodes to learn node representations in dynamic graph structures. This is particularly useful for traffic prediction, where the traffic network is inherently dynamic, with varying traffic conditions, sensor readings, and temporal fluctuations. By integrating spatio-temporal features, DST-GCNs not only model the spatial correlations between traffic nodes at a given time but also adapt to the evolving nature of traffic flow over time, enabling accurate short-term and long-term traffic predictions. Recent studies have demonstrated the effectiveness of DST-GCNs in capturing complex interactions between spatial and temporal factors, outperforming traditional traffic forecasting methods that fail to account for dynamic changes in the network.

The Temporal Graph Convolutional Network (T-GCN) [6] has emerged as a powerful approach for traffic prediction by leveraging the temporal dependencies in dynamic traffic networks. T-GCN combines the strengths of graph convolutional networks (GCNs) [3] with temporal modeling, effectively capturing both spatial correlations between traffic nodes (e.g., road segments or intersections) and their time-varying behaviors. Unlike traditional methods that treat traffic flow as static or independent at each time step, T-GCN dynamically learns the temporal evolution of traffic patterns by applying graph convolutions along with recurrent layers or temporal convolutions to model long-range dependencies. This allows T-GCN to handle irregularities in traffic data, such as peak hours or unexpected events, by learning from historical traffic conditions. Several studies have shown that T-GCN significantly improves traffic forecasting accuracy, particularly in complex urban environments, by providing a more holistic representation of traffic dynamics that traditional methods cannot capture.

Towards Dynamic Spatial-Temporal Graph Learning: A Decoupled Perspective [5] presents an innovative approach to modeling dynamic graph structures by separately addressing the spatial and temporal aspects of graph learning. This decoupled perspective aims to overcome the limitations of traditional models that treat spatial and temporal dependencies as a single, unified entity. By decoupling the spatial and temporal components, the approach allows for more flexible and efficient learning of dynamic graphs, particularly in scenarios where spatial and temporal dynamics evolve at different rates. This method enhances the model’s ability to capture the complex interactions within dynamic systems, such as traffic flow or social networks, where spatial patterns (e.g., geographical proximity) and temporal patterns (e.g., time-dependent behavior) evolve independently. The proposed framework has been shown to improve prediction accuracy in tasks involving dynamic graph data by providing a more modular and interpretable approach to learning these dual dependencies, offering new insights into handling real-time dynamic systems with greater precision.

3 PROBLEM FORMULATION

In the context of traffic modeling, we represent the traffic network as a sequence of temporal graphs $\mathcal{G}^{(t)} = (V, E, \mathbf{A})$, where V is the set of nodes, E is the set of edges, and $\mathbf{A} \in \mathbb{R}^{|V| \times |V|}$ is the adjacency matrix. Each graph captures the state of the traffic network at a specific time step t , and we have T instances of such graphs representing the temporal evolution of the network. Each node $v_i \in V$ is associated with a feature vector $\mathbf{x}_i^{(t)} \in \mathbb{R}^F$, where F represents the number of features encapsulating traffic-related characteristics such as speed, flow, and density.

For each node $v_i^{(t+1)}$ in the current graph $\mathcal{G}^{(t+1)}$, we compute the covariance matrix $\mathbf{C}_i^{(t+1)} \in \mathbb{R}^{F \times F}$ with respect to the features of all nodes $v_j^{(t)} \in \mathcal{G}^{(t)}$. Mathematically, this is expressed as:

$$\mathbf{C}_i^{(t+1)} = \frac{1}{|V^{(t)}|} \sum_{j \in V^{(t)}} \left(\mathbf{x}_j^{(t)} - \bar{\mathbf{x}}_i^{(t+1)} \right) \left(\mathbf{x}_j^{(t)} - \bar{\mathbf{x}}_i^{(t+1)} \right)^\top,$$

where $\bar{\mathbf{x}}_i^{(t+1)}$ is the mean feature vector for node $v_i^{(t+1)}$ across all time steps t . Using $\mathbf{C}_i^{(t+1)}$, eigenvalue decomposition is performed:

$$\mathbf{C}_i^{(t+1)} \mathbf{u}_k = \lambda_k \mathbf{u}_k, \quad k = 1, 2, \dots, F,$$

83 where λ_k are the eigenvalues (sorted in descending order) and \mathbf{u}_k are the corresponding eigenvectors.
 84 The K largest eigenvalues and their eigenvectors are selected, and each node's feature vector is
 85 projected onto these eigenvectors:

$$\mathbf{p}_i^{(t+1)} = \sum_{k=1}^K (\mathbf{x}_i^{(t+1)} \cdot \mathbf{u}_k) \mathbf{u}_k,$$

86 where $\mathbf{p}_i^{(t+1)}$ represents the projected features for node $v_i^{(t+1)}$. Similarly, projections $\mathbf{p}_j^{(t)}$ are
 87 computed for all nodes $v_j^{(t)} \in \mathcal{G}^{(t)}$. The similarity between the current node and previous nodes is
 88 measured using cosine similarity:

$$\text{sim}(\mathbf{p}_i^{(t+1)}, \mathbf{p}_j^{(t)}) = \frac{\mathbf{p}_i^{(t+1)} \cdot \mathbf{p}_j^{(t)}}{\|\mathbf{p}_i^{(t+1)}\| \|\mathbf{p}_j^{(t)}\|}.$$

89 Nodes $v_j^{(t)}$ with similarity greater than a threshold τ are selected, forming the set $S(v_i)$:

$$S(v_i) = \{v_j^{(t)} \in \mathcal{G}^{(t)} : \text{sim}(\mathbf{p}_i^{(t+1)}, \mathbf{p}_j^{(t)}) > \tau\}.$$

90 This set of selected nodes is used to define spatial weights $W_{i,j}$, which will further influence the
 91 aggregation of features. The final weights and feature updates will be detailed in subsequent sections.
 92 Repeat the process for all nodes $v \in V$ in G_{t+1} . For each node v , identify its corresponding set of
 93 relevant nodes $S(v)$ from the previous graph G_t . Refer to 1 algorithm. Here's the enhanced problem
 94 statement incorporating the weights learning based on attention mechanism section into the earlier
 95 description:

96 3.1 Weights Learning Based on Attention Mechanism

97 For each node $v_i \in V$ in G_{t+1} , compute weights W_{ij} for all nodes $v_j \in S(v_i)$ (i.e., the selected
 98 relevant nodes from G_t) using an **attention mechanism**. The attention weights are computed as
 99 follows:

100 3.2 Spatial Correlation via Attention

101 The spatial correlation between the features of node v_i in G_{t+1} and node v_j in G_t is given by:

$$\rho_{ij}^{(k)} = \text{LeakyReLU}(\mathbf{a}^\top \cdot \text{CONCAT}(\mathbf{U}\mathbf{x}_i, \mathbf{U}\mathbf{x}_j))$$

102 where: - \mathbf{U} is a transformation matrix applied to the feature vectors, - $\text{CONCAT}(\cdot, \cdot)$ denotes the
 103 concatenation of transformed feature vectors, - \mathbf{a} is a learnable attention vector, - LeakyReLU is the
 104 activation function applied to the attention computation.

105 3.3 Weight Coefficient Normalization

106 Normalize the attention scores using the softmax function:

$$W_{ij}^{(k)} = \frac{\exp(\rho_{ij}^{(k)})}{\sum_{u \in S(v_i)} \exp(\rho_{iu}^{(k)})}$$

107 This normalization ensures that the weights sum to 1 across the selected nodes. 2 algorithm

Algorithm 1 Feature Embedding and Node Selection for Traffic Graphs

Input: Sequence of temporal graphs $\{\mathcal{G}^{(t)} = (V, E, \mathbf{A})\}_{t=0}^{T-1}$, Node features $\mathbf{x}_i^{(t)} \in \mathbb{R}^F$, Threshold τ , Number of top eigenvectors K .

Output: Set of selected nodes $S(v_i)$ for each node v_i in $\mathcal{G}^{(t+1)}$.

for each node $v_i \in V^{(t+1)}$ **do**

 Initialize $S(v_i) = \emptyset$

 Compute the mean feature vector:

$$\bar{\mathbf{x}}_i^{(t+1)} = \frac{1}{|V^{(t)}|} \sum_{j \in V^{(t)}} \mathbf{x}_j^{(t)}$$

 Compute the covariance matrix:

$$\mathbf{C}_i^{(t+1)} = \frac{1}{|V^{(t)}|} \sum_{j \in V^{(t)}} \left(\mathbf{x}_j^{(t)} - \bar{\mathbf{x}}_i^{(t+1)} \right) \left(\mathbf{x}_j^{(t)} - \bar{\mathbf{x}}_i^{(t+1)} \right)^\top$$

 Perform eigenvalue decomposition:

$$\mathbf{C}_i^{(t+1)} \mathbf{u}_k = \lambda_k \mathbf{u}_k, \quad k = 1, 2, \dots, F$$

 Sort eigenvalues λ_k in descending order and select the top K eigenvectors $\{\mathbf{u}_k\}_{k=1}^K$

 Project current node features onto eigenvectors:

$$\mathbf{p}_i^{(t+1)} = \sum_{k=1}^K (\mathbf{x}_i^{(t+1)} \cdot \mathbf{u}_k) \mathbf{u}_k$$

for each node $v_j \in V^{(t)}$ **do**

 Project features of v_j onto eigenvectors:

$$\mathbf{p}_j^{(t)} = \sum_{k=1}^K (\mathbf{x}_j^{(t)} \cdot \mathbf{u}_k) \mathbf{u}_k$$

 Compute cosine similarity:

$$\text{sim}(\mathbf{p}_i^{(t+1)}, \mathbf{p}_j^{(t)}) = \frac{\mathbf{p}_i^{(t+1)} \cdot \mathbf{p}_j^{(t)}}{\|\mathbf{p}_i^{(t+1)}\| \|\mathbf{p}_j^{(t)}\|}$$

if $\text{sim} > \tau$ **then**

 Add v_j to $S(v_i)$

end if

end for

end for

return $\{S(v_i) : v_i \in V^{(t+1)}\}$

108 3.4 Global Spatial Weight Matrix

109 Combine weights across all feature dimensions k to construct a global weight matrix:

$$W = \sum_{k=1}^F W^{(k)}$$

110 3.5 Multi-head Attention for Robustness

111 In practice, use multi-head attention to improve the stability and robustness of weight learning. Let M
 112 denote the number of attention heads. The final weight coefficients are averaged across all attention
 113 heads:

$$W_{ij} = \frac{1}{M} \sum_{m=1}^M W_{ij,m}$$

114 where $W_{ij,m}$ is the weight computed by the m -th attention module. 3 algorithm

115 Given consecutive input feature matrices $\mathbf{X}^{(0)}, \mathbf{X}^{(1)}, \dots, \mathbf{X}^{(S-1)}$, where $\mathbf{X}^{(t)}$ represents the feature
 116 vectors of all nodes at time t , and the spatial weights \mathbf{W} obtained using the attention mechanism,
 117 we aim to generate embeddings $\tilde{\mathbf{X}}^{(S)} = \{\tilde{\mathbf{x}}_i^{(S)} \mid \forall i \in V\}$ that summarize the temporal and spatial
 118 relationships. The process begins with feature aggregation at each time step t . For each node $v_i \in V$,
 119 features from the selected nodes $v_j \in S(v_i)$, as determined in the earlier steps, are aggregated using
 120 the spatial weight matrix \mathbf{W} . The aggregated feature vector \mathbf{h}_i is computed as:

$$\mathbf{h}_i = \sum_{j \in S(v_i)} W_{i,j} \mathbf{x}_j^{(t)},$$

121 where $W_{i,j}$ represents the spatial weight between node v_i and a selected node $v_j \in S(v_i)$, and $\mathbf{x}_j^{(t)}$
 122 is the feature vector of node v_j at time t . 4 algorithm. For each node v_i , the current feature $\mathbf{x}_i^{(t)}$ is
 123 concatenated with the aggregated features \mathbf{h}_i , resulting in the concatenated feature vector:

$$\mathbf{h}_i^{\text{cat}} = \text{CONCAT}(\mathbf{x}_i^{(t)}, \mathbf{h}_i).$$

124 The concatenated feature vector $\mathbf{h}_i^{\text{cat}} \in \mathbb{R}^{2F}$ (since it combines $\mathbf{x}_i^{(t)} \in \mathbb{R}^F$ and $\mathbf{h}_i \in \mathbb{R}^F$) is then
 125 passed through a fully connected layer to produce a transformed feature representation for node v_i .
 126 This transformation is given by:

$$\mathbf{h}_i^{(t)} = \sigma(\mathbf{h}_i^{\text{cat}} \mathbf{U}_t + \mathbf{b}_t),$$

127 where $\mathbf{U}_t \in \mathbb{R}^{2F \times F'}$ is a learnable weight matrix, $\mathbf{b}_t \in \mathbb{R}^{F'}$ is a learnable bias vector, and $\sigma(\cdot)$
 128 is an activation function (e.g., ReLU). The transformed feature vector $\mathbf{h}_i^{(t)} \in \mathbb{R}^{F'}$ represents the
 129 embedding of node v_i at time t . This process is repeated for S consecutive time steps, 5 algorithm
 130 and the final temporal-spatial embedding $\tilde{\mathbf{x}}_i^{(S)}$ for node v_i is computed by summing the transformed
 131 embeddings over all S time steps:

$$\tilde{\mathbf{x}}_i^{(S)} = \sum_{t=0}^{S-1} \mathbf{h}_i^{(t)}.$$

132 Refer to 6 algorithm. The final embedding matrix $\tilde{\mathbf{X}}^{(S)} \in \mathbb{R}^{N \times F'}$ is obtained by aggregating the
 133 embeddings of all nodes in the graph, where each row corresponds to the embedding $\tilde{\mathbf{x}}_i^{(S)}$ of a node
 134 $v_i \in V$. By restricting the aggregation to nodes in $S(v_i)$, this approach incorporates only the most
 135 relevant spatial and temporal relationships for each node, as determined by the cosine similarity
 136 thresholding in the earlier steps. This ensures that the final embeddings capture meaningful patterns

137 in both space and time, making them suitable for downstream tasks such as prediction, clustering, or
 138 anomaly detection.

Algorithm 2 Part 1: Attention Weight Computation

- 1: **Input:** Feature matrices $\mathbf{X}^{(0)}, \mathbf{X}^{(1)}, \dots, \mathbf{X}^{(S-1)}$ for S time steps. Each $\mathbf{X}^{(t)} \in \mathbb{R}^{N \times F}$ represents the node features at time step t .
- 2: **Output:** Attention weights matrix W_{ij} .
- 3: **for** $t = 0$ **to** $S - 1$ **do**
- 4: **For each time step** t :
- 5: **for** each node $v_i \in V$ **do**
- 6: **Step A:** Compute attention weights for each selected node $v_j \in S(v_i)$:

$$W_{ij}^{(k)} = \frac{\exp(\rho_{ij}^{(k)})}{\sum_{u \in S(v_i)} \exp(\rho_{iu}^{(k)})}$$

where $\rho_{ij}^{(k)}$ is computed as:

$$\rho_{ij}^{(k)} = \text{LeakyReLU} \left(\mathbf{a}^\top \cdot \text{CONCAT}(\mathbf{U}\mathbf{x}_i^{(t)}, \mathbf{U}\mathbf{x}_j^{(t)}) \right)$$

- 7: **end for**
 - 8: **end for**
-

Algorithm 3 Part 2: Weight Normalization and Multi-head Attention

- 1: **Input:** Attention weights matrix W_{ij} .
- 2: **Output:** Normalized weights matrix W_{ij} and multi-head attention weights.
- 3: **for** $t = 0$ **to** $S - 1$ **do**
- 4: **For each time step** t :
- 5: **for** each node $v_i \in V$ **do**
- 6: **Step B:** Normalize the attention scores using the softmax function:

$$W_{ij}^{(k)} = \frac{\exp(\rho_{ij}^{(k)})}{\sum_{u \in S(v_i)} \exp(\rho_{iu}^{(k)})}$$

- 7: **end for**
- 8: **end for**
- 9: **Step C:** Apply Multi-head Attention to improve robustness:

$$W_{ij} = \frac{1}{M} \sum_{m=1}^M W_{ij,m}$$

where M is the number of attention heads, and $W_{ij,m}$ is the weight computed by the m -th attention head.

139 4 PRELIMINARY RESULTS

140 In our preliminary experiments, we explored a modified version of the graph convolutional method
 141 which includes the temporal covariance of the nodes into the convolutional operation. Equation 1
 142 shows the original formulation of graph convolutions, where the output $H^{(l+1)}$ is computed as a
 143 function of the input feature matrix $H^{(l)}$, the weight matrix $W^{(l)}$, and the normalized Laplacian
 144 matrix L_{sym} .

$$H^{(l+1)} = \sigma(L_{\text{sym}} H^{(l)} W^{(l)}) \quad (1)$$

Algorithm 4 Part 3: Feature Aggregation

- 1: **Input:** Feature matrix $\mathbf{X}^{(t)}$ and normalized attention weights W_{ij} .
- 2: **Output:** Aggregated feature vector $\mathbf{h}_i^{(t)}$.
- 3: **for** $t = 0$ **to** $S - 1$ **do**
- 4: **For each time step** t :
- 5: **for** each node $v_i \in V$ **do**
- 6: **Step D:** Aggregate features from selected neighbors $v_j \in S(v_i)$ using the attention weights:

$$\mathbf{h}_i^{(t)} = \sum_{j \in S(v_i)} W_{ij} \mathbf{x}_j^{(t)}$$

- 7: **end for**
 - 8: **end for**
-

Algorithm 5 Part 4: Feature Concatenation and Transformation

- 1: **Input:** Original feature vector $\mathbf{x}_i^{(t)}$ and aggregated feature $\mathbf{h}_i^{(t)}$.
- 2: **Output:** Transformed feature vector $\mathbf{h}_i^{(t)}$.
- 3: **for** $t = 0$ **to** $S - 1$ **do**
- 4: **For each time step** t :
- 5: **for** each node $v_i \in V$ **do**
- 6: **Step E:** Concatenate the original feature with the aggregated feature:

$$\mathbf{h}_i^{\text{cat}} = \text{CONCAT}(\mathbf{x}_i^{(t)}, \mathbf{h}_i^{(t)})$$

- 7: **Step F:** Transform the concatenated feature via a fully connected layer:

$$\mathbf{h}_i^{(t)} = \sigma(\mathbf{h}_i^{\text{cat}} \mathbf{U}_t + \mathbf{b}_t)$$

where $\mathbf{U}_t \in \mathbb{R}^{2F \times F'}$ and $\mathbf{b}_t \in \mathbb{R}^{F'}$ are learnable parameters, and $\sigma(\cdot)$ is an activation function (e.g., ReLU).

- 8: **end for**
 - 9: **end for**
-

Algorithm 6 Part 5: Temporal-Spatial Embedding Generation

- 1: **Input:** Transformed feature vectors $\mathbf{h}_i^{(t)}$ for all time steps.
- 2: **Output:** Final temporal-spatial embedding $\tilde{\mathbf{x}}_i^{(S)}$.
- 3: **for** each node $v_i \in V$ **do**
- 4: **Step G:** Generate final temporal-spatial embedding by aggregating the transformed embeddings over all time steps:

$$\tilde{\mathbf{x}}_i^{(S)} = \sum_{t=0}^{S-1} \mathbf{h}_i^{(t)}$$

- 5: **end for**
- 6: **Step H:** Construct the final embedding matrix:

$$\tilde{\mathbf{X}}^{(S)} = [\tilde{\mathbf{x}}_1^{(S)}, \tilde{\mathbf{x}}_2^{(S)}, \dots, \tilde{\mathbf{x}}_N^{(S)}]$$

145 To further enhance the model’s ability to capture temporal dependencies between nodes, we incorpo-
 146 rated covariance-based embeddings, as proposed in the recent work by [1]. Equation 2 illustrates
 147 the updated approach, where we apply element-wise masks M_l and M_c to the Laplacian matrix
 148 L_{sym} and the covariance matrix C , respectively. These masks selectively weight or zero out certain
 149 contributions, allowing the model to learn more robust representations.

$$H^{(l+1)} = \sigma \left(\left(\frac{M_l \odot L_{sym} + M_c \odot C}{M_l \odot L_{sym} + M_c \odot C} \right) H^{(l)} W^{(l)} \right) \quad (2)$$

150 By incorporating the covariance matrix C , which models the temporal dependencies between nodes,
 151 we aim to improve the model’s performance on tasks that require understanding the dynamic re-
 152 lationships within the data. In the following, we present our experimental results comparing the
 153 performance of the original graph convolutional method and the modified version with the covariance-
 154 based embeddings. This will allow us to assess the impact of the proposed changes and their potential
 155 to enhance the model’s ability to capture complex spatio-temporal patterns in the data.

156 4.1 About the data

- 157 1. **SZ-taxi Dataset:** This dataset contains taxi trajectory data from Shenzhen, collected between
 158 January 1 and January 31, 2015. The study focuses on 156 major roads in the Luohu District.
 159 The dataset comprises two main components: an adjacency matrix and a feature matrix. The
 160 156×156 adjacency matrix represents the spatial relationships between roads, where each
 161 row corresponds to a road, and the values of the matrix indicate the connectivity between
 162 them. The feature matrix captures changes in traffic speed over time for each road. Each
 163 row corresponds to a road, while each column represents the aggregated traffic speed in
 164 15-minute intervals.
- 165 2. **Los-loop Dataset:** This dataset was gathered from real-time loop detectors on highways in
 166 Los Angeles County between March 1 and March 7, 2012. It includes data from 207 sensors
 167 and records their traffic speed, aggregated every 5 minutes. Similarly to the SZ-taxi dataset,
 168 it comprises an adjacency matrix and a feature matrix. The adjacency matrix is based on
 169 the distances between the sensors within the traffic network. Since the dataset contained
 170 missing values, a linear interpolation method was applied to fill the gaps.

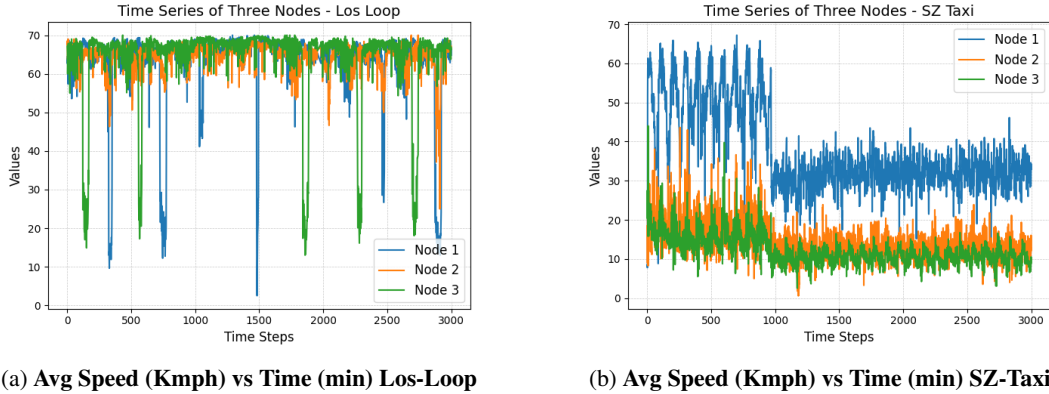


Figure 1: Time Series data of first three nodes

171 Both datasets exhibit distinct characteristics. The Los-Loop dataset demonstrates relatively stable
 172 patterns with minimal distributional shifts, as evident from the data of the first three nodes shown in
 173 Figure 1a. In contrast, the SZ-Taxi dataset shows noticeable distributional shifts, with a clear shift in
 174 the mean over time, as illustrated in Figure 1b.

175 4.2 Experiments

176 We leveraged the code bases of TGCN [6] and Spatio-Temporal Covariance Neural Networks [1],
 177 for reference and combined the batch covariance calculation from the latter into the former. For the

178 experiments, we perform forecasting the average speed for 15min, 30min, 45min, 60min respectively
 179 for both the models and evaluate the model performance on certain metrics to compare the model
 180 performances. We show that the cVTGCN model which includes temporal covariance embeddings
 181 performs better than its vanilla variant.

182 4.2.1 Model Design

183 We employ the TGCN model as the foundation, enhancing it with an additional overlay layer that
 184 incorporates covariance calculations and masking operations. The key hyperparameters of the TGCN
 185 model include the learning rate, batch size, number of training epochs, and the number of hidden
 186 layers. For this experiment, these hyperparameters were fine-tuned and set to the following values: a
 187 learning rate of 0.001, a batch size of 64, and 100 training epochs.

188 During the training process, 80% of the overall dataset is used as the training set, while the remaining
 189 20% is reserved for testing. The TGCN model is trained using the Adam optimizer to ensure efficient
 190 and stable convergence.

191 To assess the prediction performance of the TGCN model, we evaluate the difference between the
 192 actual traffic information Y_t and the predicted values \hat{Y}_t using three metrics. These metrics provide a
 193 comprehensive analysis of the model's accuracy and robustness.

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (Y_t - \hat{Y}_t)^2}$$

$$MAE = \frac{1}{n} \sum_{i=1}^n |Y_t - \hat{Y}_t|$$

$$Accuracy = 1 - \frac{\|Y - \hat{Y}\|_F}{\|Y\|_F}$$

196 4.2.2 Loss Function

197 During training, the objective is to minimize the discrepancy between the actual and predicted traffic
 198 speeds, denoted as Y_t and \hat{Y}_t , respectively. The loss function for the T-GCN model is defined as:

$$\text{Loss} = \|Y_t - \hat{Y}_t\|^2 + \lambda L_{\text{reg}} + \mu L_{\text{cov}} \quad (3)$$

199 The first term minimizes the prediction error, ensuring the model accurately captures traffic dynamics.
 200 The second term, L_{reg} , is an L_2 regularization term to mitigate overfitting, with λ as its weight.
 201 The third term, L_{cov} , introduces covariance-based regularization to capture temporal dependencies
 202 effectively, weighted by the hyperparameter μ . The covariance-based regularization term, L_{cov} , is
 203 defined as the Frobenius norm of the difference between the predicted and observed covariance
 204 matrices:

$$L_{\text{cov}} = \|\mathbf{C}_{\text{pred}} - \mathbf{C}_{\text{obs}}\|_F^2$$

205 where:

- 206 • \mathbf{C}_{pred} is the predicted covariance matrix of traffic speeds.
- 207 • \mathbf{C}_{obs} is the observed covariance matrix of traffic speeds.
- 208 • $\|\cdot\|_F$ represents the Frobenius norm.


209 By incorporating L_{cov} , the model aligns predictions with observed temporal covariance patterns,
 210 enhancing its ability to generalize and capture complex traffic dynamics.

4.3 Analysis of results


From Table 1, it is evident that cVTGCN outperforms significantly on the SZ-Taxi dataset, while the improvement in accuracy for the Los-Loop dataset is comparatively marginal. This highlights cVTGCN’s ability to effectively capture distributional shifts, as the incorporation of covariance embeddings within the convolutional operations enables the model to identify more nuanced patterns in the time series data.

Table 1: Validation Metrics at Different Time Intervals

T	Metric	Los-loop		SZ-taxi	
		TGCN	cVTGCN (Ours)	TGCN	cVTGCN (Ours)
15min	RMSE	5.3861	5.2665	5.9953	4.2328
	MAE	3.5388	3.2937	4.4205	2.9625
	Accuracy	0.9083	0.9103	0.5822	0.7050
30min	RMSE	7.6295	6.3116	6.0514	4.3108
	MAE	5.0990	3.8155	4.4882	3.0719
	Accuracy	0.8701	0.8925	0.5783	0.6996
45min	RMSE	8.3162	7.1820	6.0921	4.3536
	MAE	5.4171	4.3116	4.5216	3.1274
	Accuracy	0.8583	0.8777	0.5753	0.6965
60min	RMSE	8.8854	7.9028	6.1159	4.4775
	MAE	5.7577	4.7294	4.5468	3.2194
	Accuracy	0.8486	0.8653	0.5736	0.6878



(a) Accuracy



(b) RMSE

Figure 2: Comparison of methods with SZ-Taxi

From Figure 3, it is evident that the blue line representing cVTGCN aligns more closely with the mean compared to TGCN, indicating its improved ability to capture patterns. Similarly, in Figure 4, the blue line for cVTGCN demonstrates a more pronounced shift toward the mean, further validating its superior capability in capturing distributional shifts.

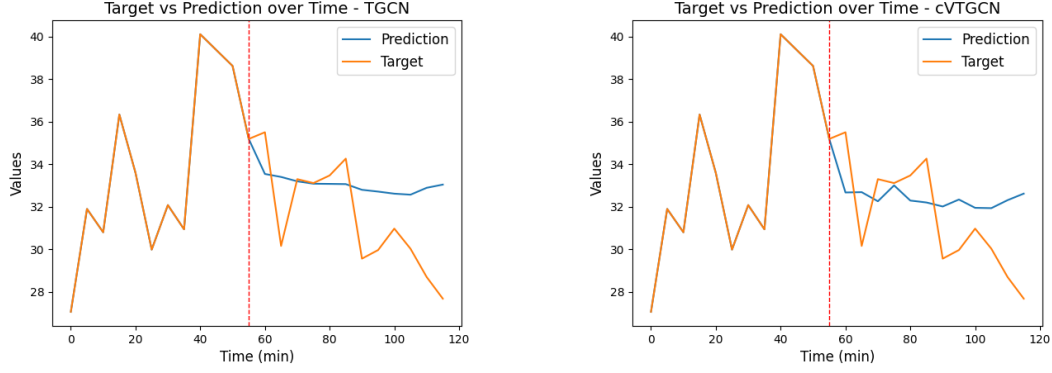


Figure 3: Target vs Prediction (Los-Loop): Target 12 timesteps - Output 12 timesteps

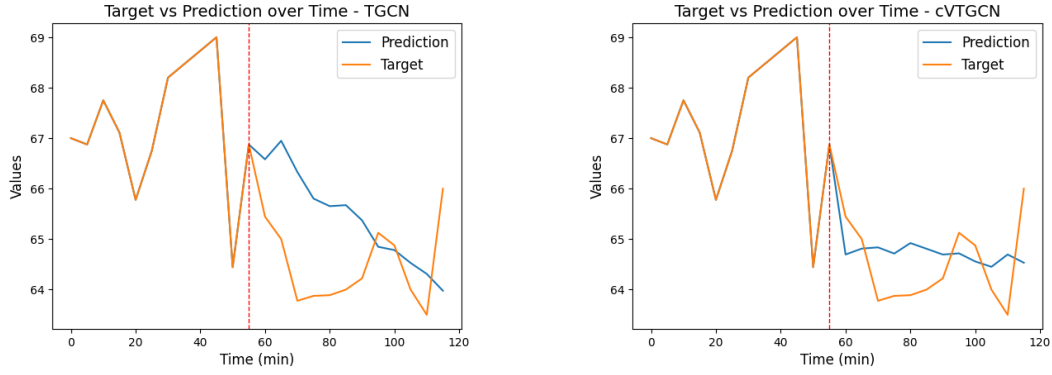


Figure 4: Input vs Prediction (SZ-Taxi): Target 12 timesteps - Output 12 timesteps

5 CONTRIBUTION STATEMENT

The contributions to this work are as follows:

- **Ideation:** Vamsi and Vikas
- **Presentation Slides:** Vamsi and Vikas
- **Literature Review:** Vamsi and Vikas
- **Code Implementation / Results:** Vamsi
- **Report Writing:** Vikas

References

- [1] Andrea Cavallo, Mohammad Sabbaghi, and Elvin Isufi. Spatiotemporal covariance neural networks. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pages 18–34. Springer, 2024.
- [2] Will Hamilton, Zhitao Ying, and Jure Leskovec. Inductive representation learning on large graphs. *Advances in neural information processing systems*, 30, 2017.
- [3] Thomas N Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. *arXiv preprint arXiv:1609.02907*, 2016.

- 236 [4] Tao Liu, Aimin Jiang, Jia Zhou, Min Li, and Hon Keung Kwan. Graphsage-based dynamic
237 spatial-temporal graph convolutional network for traffic prediction. *IEEE Transactions on*
238 *Intelligent Transportation Systems*, 24(10):11210–11224, 2023.
- 239 [5] Binwu Wang, Pengkun Wang, Yudong Zhang, Xu Wang, Zhengyang Zhou, Lei Bai, and Yang
240 Wang. Towards dynamic spatial-temporal graph learning: A decoupled perspective. *Proceedings*
241 *of the AAAI Conference on Artificial Intelligence*, 38(8):9089–9097, Mar. 2024.
- 242 [6] Ling Zhao, Yujiao Song, Chao Zhang, Yu Liu, Pu Wang, Tao Lin, Min Deng, and Haifeng Li.
243 T-gcn: A temporal graph convolutional network for traffic prediction. *IEEE transactions on*
244 *intelligent transportation systems*, 21(9):3848–3858, 2019.